



Spectrum: Speedy and Strictly-Deterministic Smart Contract Transactions for Blockchain Ledgers

Zhihao Chen, Tianji Yang, Yixiao Zheng, Zhao Zhang,
Cheqing Jin, Aoying Zhou

East China Normal University
chenzh@stu.ecnu.edu.cn



华东师范大学
EAST CHINA NORMAL UNIVERSITY



SCHOOL OF DATA
SCIENCE & ENGINEERING
数据科学与工程学院

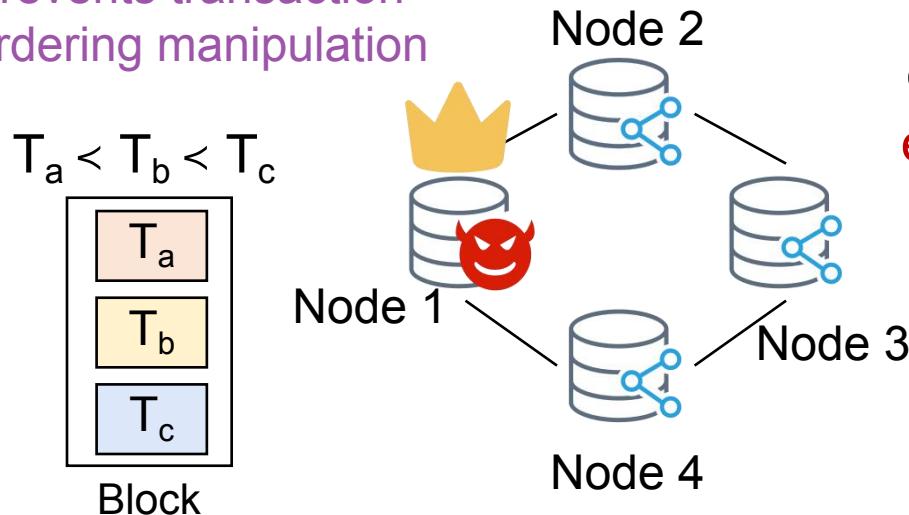
Outline

- Motivation
- Background
- Goals & Contributions
- Methodology
- Evaluation
- Conclusion & Future work

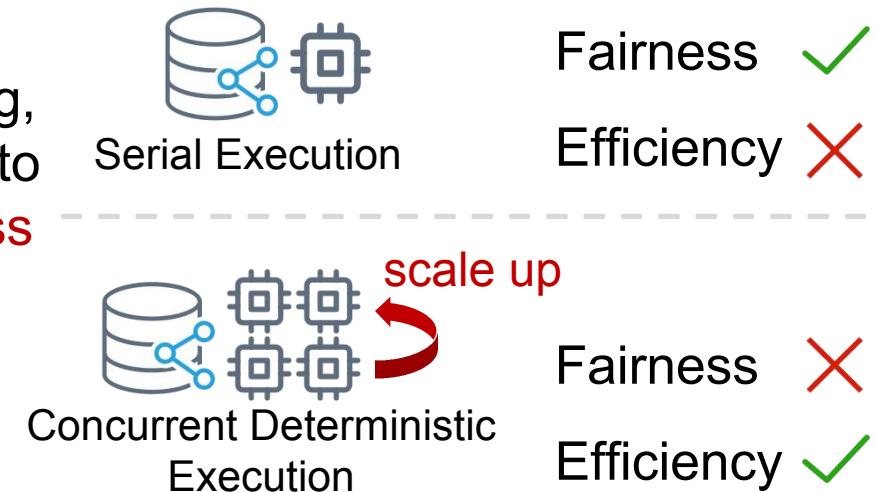
Motivation

- Unlike traditional databases, **blockchain ledgers** concern **ordering fairness**
- Existing deterministic execution schemes fail to preserve both ordering fairness and high performance

Prevents transaction ordering manipulation



Given the fair ordering, **execution** is required to preserve such fairness



Modern Byzantine consensus, e.g., Pompe [OSDI' 20]、
Themis [CCS' 23] incorporate fairness designs

Motivation

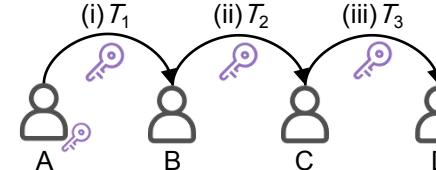
- They merely guarantee deterministic serializability, which can lead to a **deterministic (consistent across replicas)**, yet different serial order than the agreed-upon ordering

e.g., **AccessControl contract**: A caller with access permission can grant it to a specified address

```
contract AccessControl {  
    mapping(address => bool) public access;  
    function grantAccess(address to) public {  
        if (access[msg.sender] != false)  
            access[to] = access[msg.sender];  
    }  
}
```

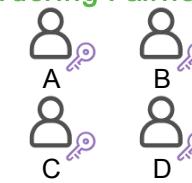
Initially, only A has permission; T_1 : A->B, T_2 : B->C, T_3 : C->D

The Fair Consensus Ordering: $T_1 \rightarrow T_2 \rightarrow T_3$



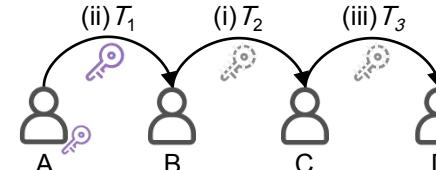
Execution Preserves Ordering Fairness

After exec.



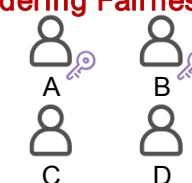
Diverged results

A Deterministic Execution Serial Order: $T_2 \rightarrow T_1 \rightarrow T_3$



Execution Disrupts Ordering Fairness

After exec.

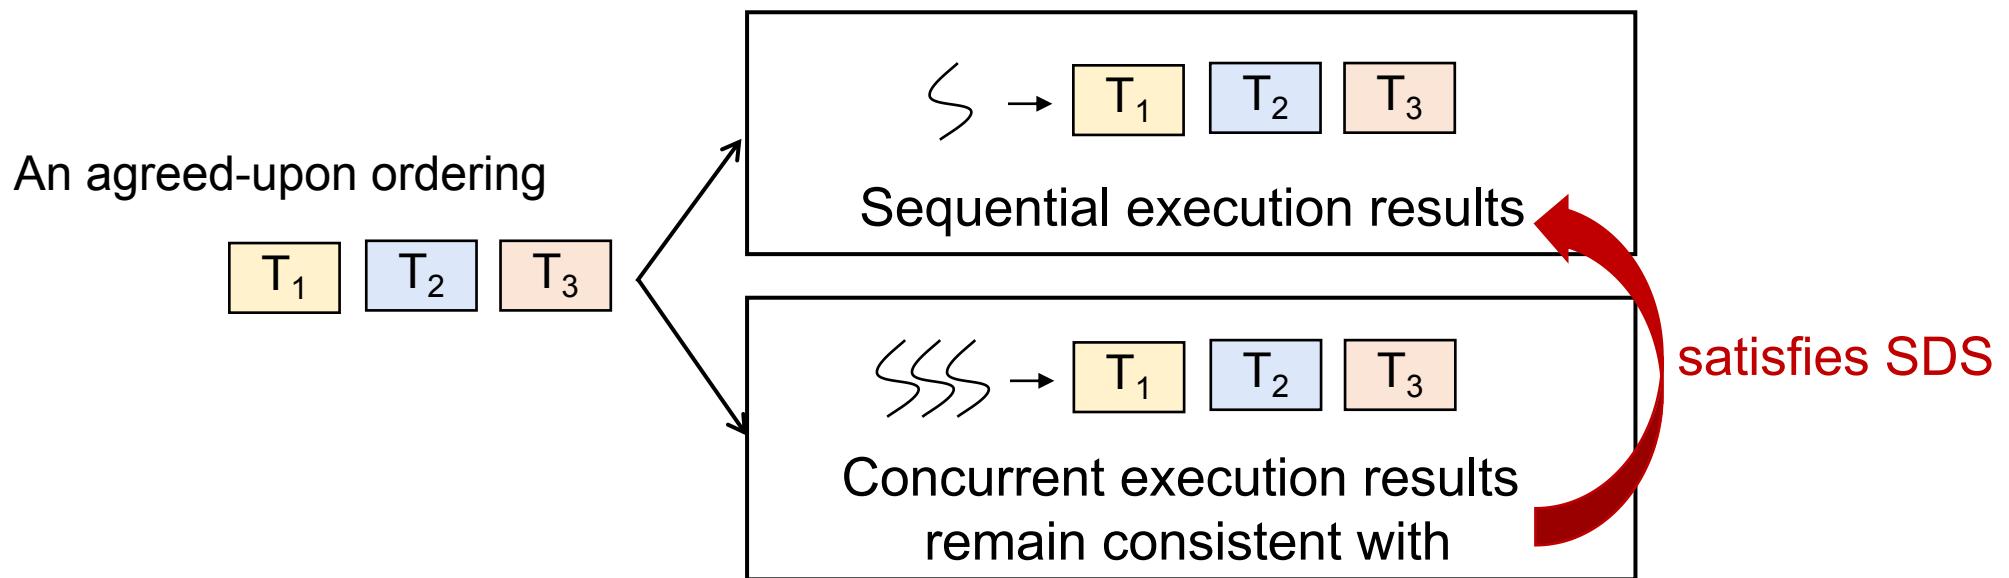


Merely Determinism Does NOT Preserve Ordering Fairness

Background

■ Problem Definition: Strictly-Deterministic Serializability (SDS)

- Given an agreed ordering of transactions, $O: \{T_1, \dots, T_n\}$, an execution schedule of transactions S satisfies **strictly-deterministic serializability** iff its effect is equivalent to the sequential execution of O , which adheres to the transactions commit order, $\{T_1, \dots, T_n\}$



Background

- For blockchain ledgers, ensuring SDS in an execution scheme preserves ordering fairness
- But most Deterministic Concurrency Control (DCC) schemes **fail to ensure SDS** when processing **smart contracts with runtime-determined accesses**

Deterministic
Concurrency
Control
(Database)

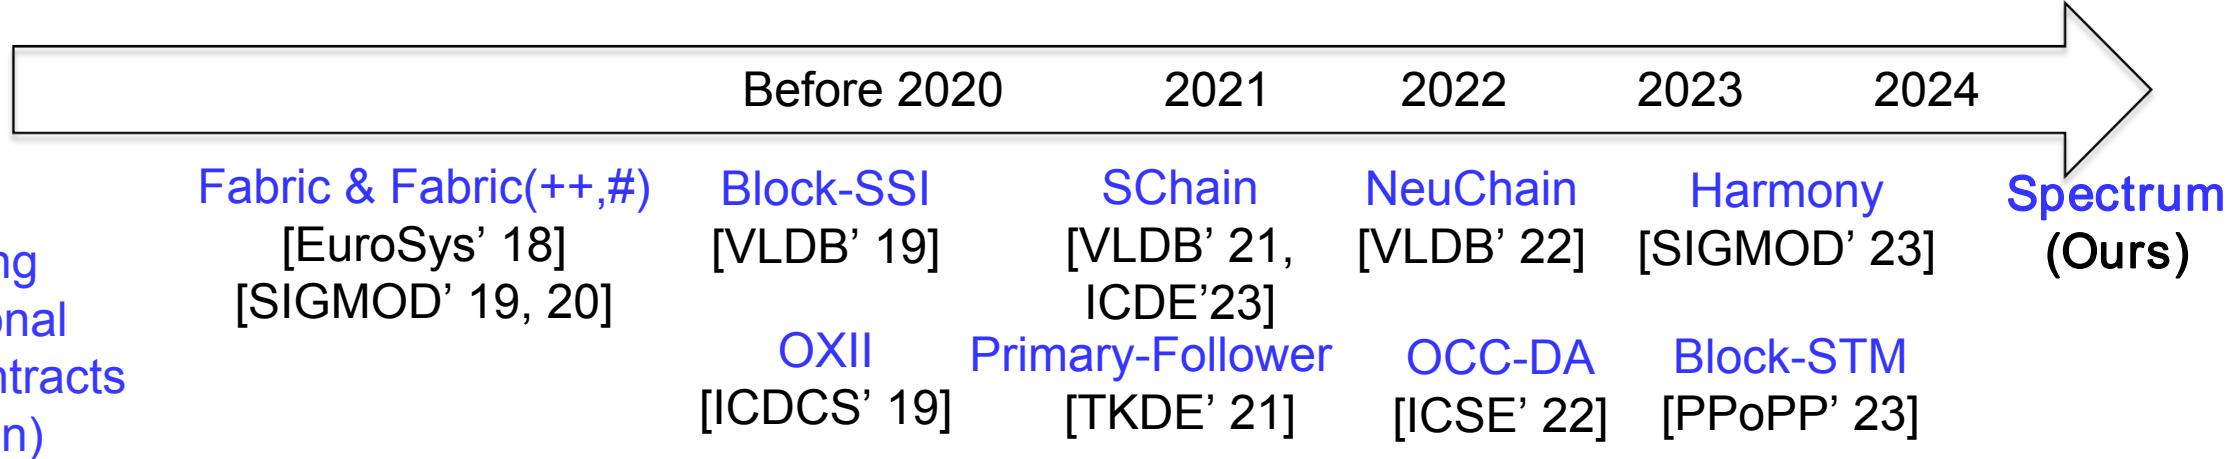
Calvin	Orthrus	Sparkle	QStore	Caracal
[SIGMOD' 12]	[SIGMOD' 16]	[DSN' 19]	[EDBT' 20]	[SOSP' 21]
Bohm	PWV	QueCC	Aria	Lotus
[VLDB' 15]	[VLDB' 17]	[Middleware' 18]	[VLDB' 20]	[VLDB' 22]

Kindly refer to the paper
for an in-depth analysis

Deterministic Execution Schemes	Execution Paradigm	No Complete R/W Sets	Strict Determinism	Scaling in Contention
HyperLedger Fabric [10]	EOV	✓	✗	✗
Fabric(++, #) [47, 49]	EOV	✓	✗	✓
NeuChain [40]	EV	✓	✗	✗
Calvin [55], PWV [22]	OE	✗	✗	✓
Bohm [23], Caracal [45]	OE	✗	✗	✓
QueCC [43]	OE	✗	✗	✓
Aria, AriaPB [38]	OE	✓	✗	✓
Sparkle [36]	OE	✓	✓	✗
Primary-Follower [29]	OE	✓	✗	✓
SSI [39], OCC-DA [24]	OE	✓	✗	✗
OXII [9], PEEP [15]	OE	✗	✗	✓
Harmony [33]	OE	✓	✗	✓
Spectrum (Ours)	OE	✓	✓	✓

* Scope of Databases

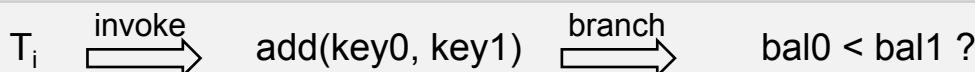
▲ Scope of Blockchains



Background

- Quasi-Turing-complete smart contracts have **runtime-determined access patterns**

```
1 pragma solidity >=0.8.2 <0.9.0;
2 contract MutableRW {
3     mapping(uint256 => uint256) public store;
4     function add(uint256 key0, uint256 key1) public {
5         uint256 bal0 = store[key0];
6         uint256 bal1 = store[key1];
7         if (bal0 < bal1) store[key0] += bal1;
8         else store[key1] += bal0;
9     }
10 }
```



Branch 1
 $\text{Snapshot}_1 \rightarrow (\text{bal0} < \text{bal1})$

T_i 's R_Set: { $\text{store}[\text{key1}]$ }

T_i 's RW_Set: { $\text{store}[\text{key0}]$ } R/W Sets

Branch 2
 $\text{Snapshot}_2 \rightarrow (\text{bal0} \geq \text{bal1})$

T_i 's R_Set: { $\text{store}[\text{key0}]$ }

T_i 's RW_Set: { $\text{store}[\text{key1}]$ }

A Transaction with mutable read/write sets

Runtime-determined nature



whose read/write sets can **vary** across different snapshots
(Mutable read/write sets)



Most DCC schemes **CAN NOT** guarantee SDS when handling mutable r/w sets

For inaccurate pre-acquisition or inherent scheme limitations (reordering, violations, etc.)

Goals & Contributions

■ Design goal

- Spectrum, a DCC scheme that ensures **both strict determinism and high performance (across diverse workloads)** for blockchain ledgers

■ Key points and contributions

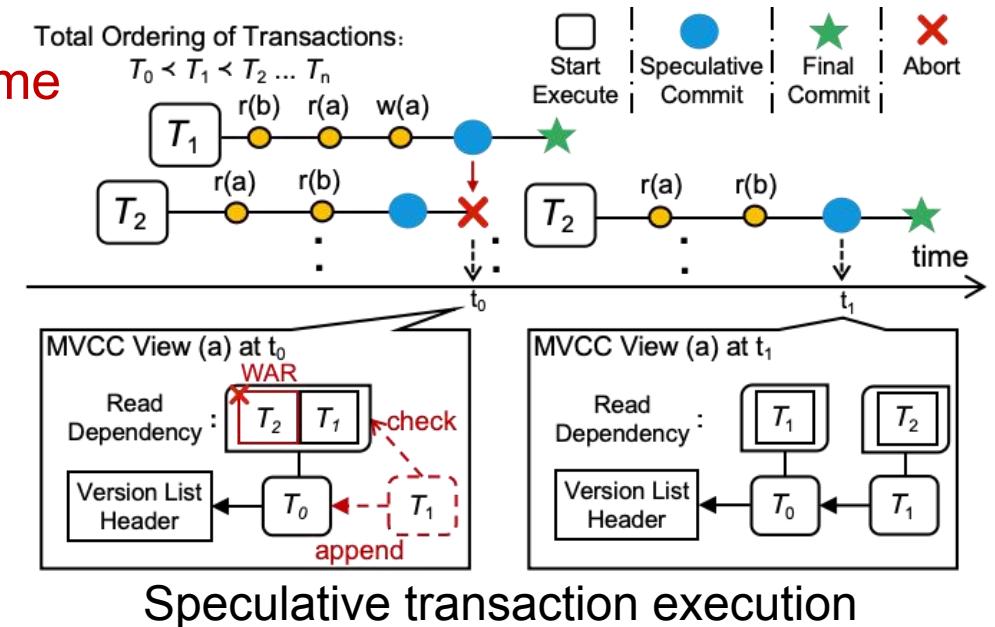
- Leverages speculative execution to **ensure SDS for concurrent smart contract transactions**
 - Proposes a partial rollback mechanism with efficient impl.
 - Designs a predictive transaction scheduling method
 - Evaluates by running EVM-Based smart contracts on YCSB, SmallBank and TPC-C alike benchmarks
- Two novel optimizations to maintain high performance under contention

Methodology

■ #1 Speculative transaction execution

- Multi-Version Concurrency Control with runtime conflict detection
- Lets each thread independently execute Txns in speculation => high inter-thread concurrency
- Transaction lifecycle: 1) Execution; 2) Speculative Commit; 3) Final Commit

- Detects and aborts **any order-violating Txns at runtime**
- **Re-executes them** with their original seq. numbers
- Upholds SDS

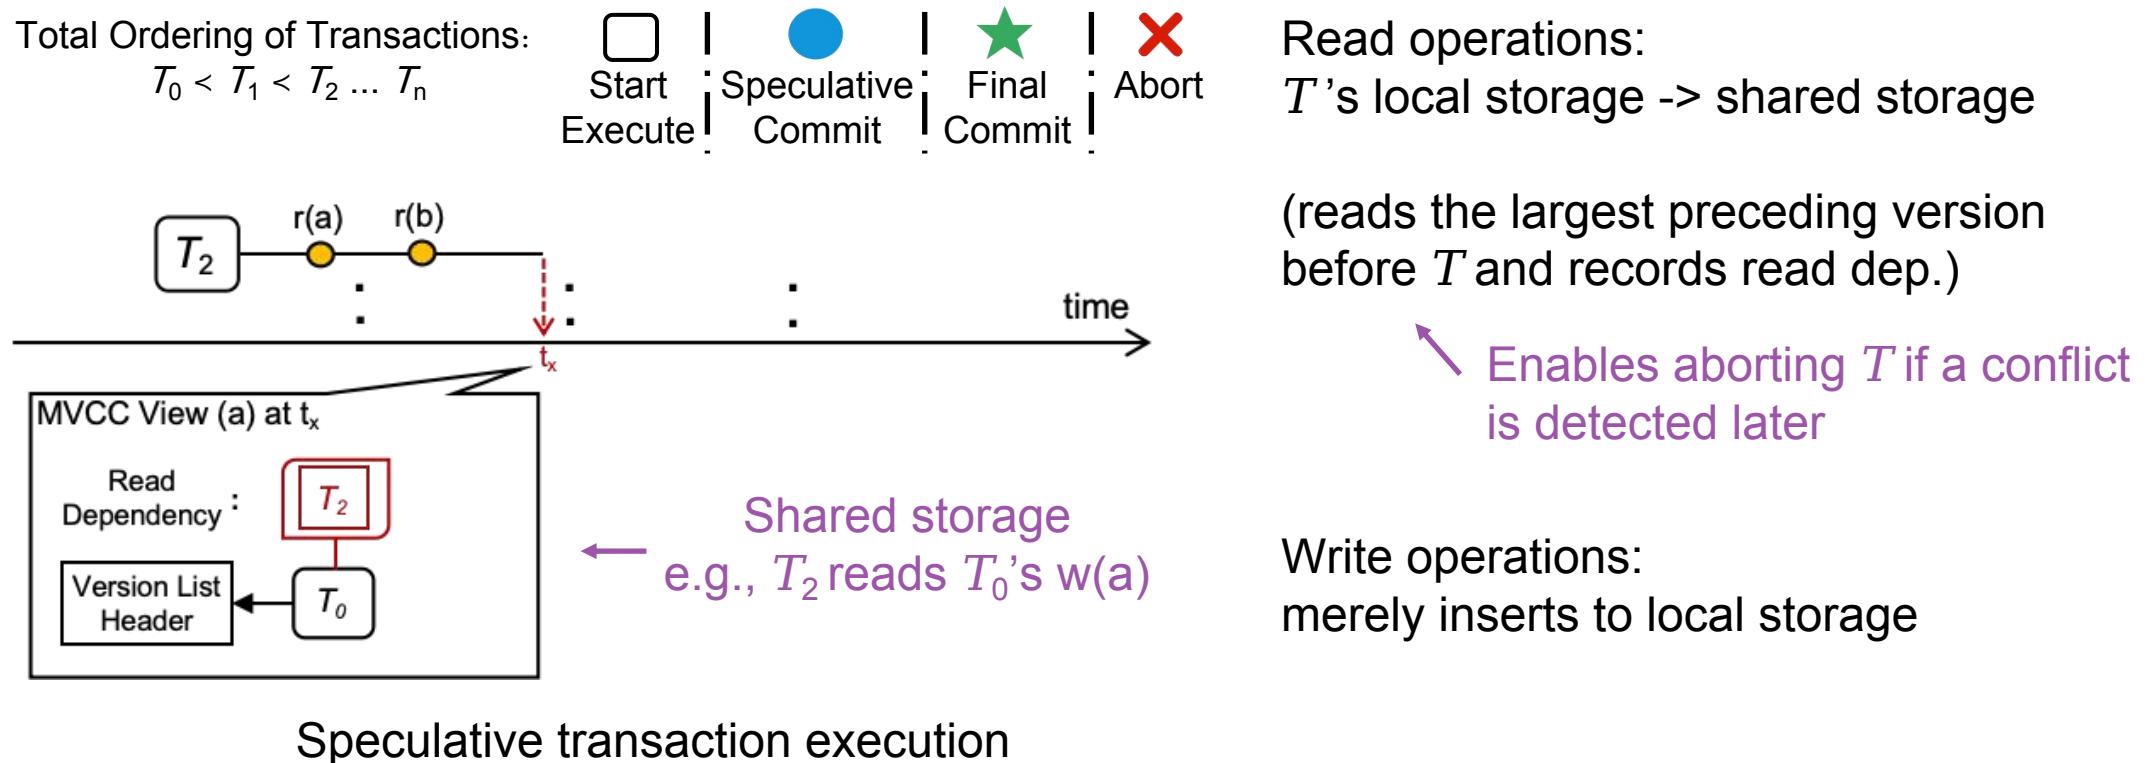


Methodology

■ #1 Speculative transaction execution

- Transaction lifecycle: 1) Execution; 2) Speculative Commit; 3) Final Commit

1) Execution: runs OPs, reads from shared storage (if not found locally), writes to T 's local storage

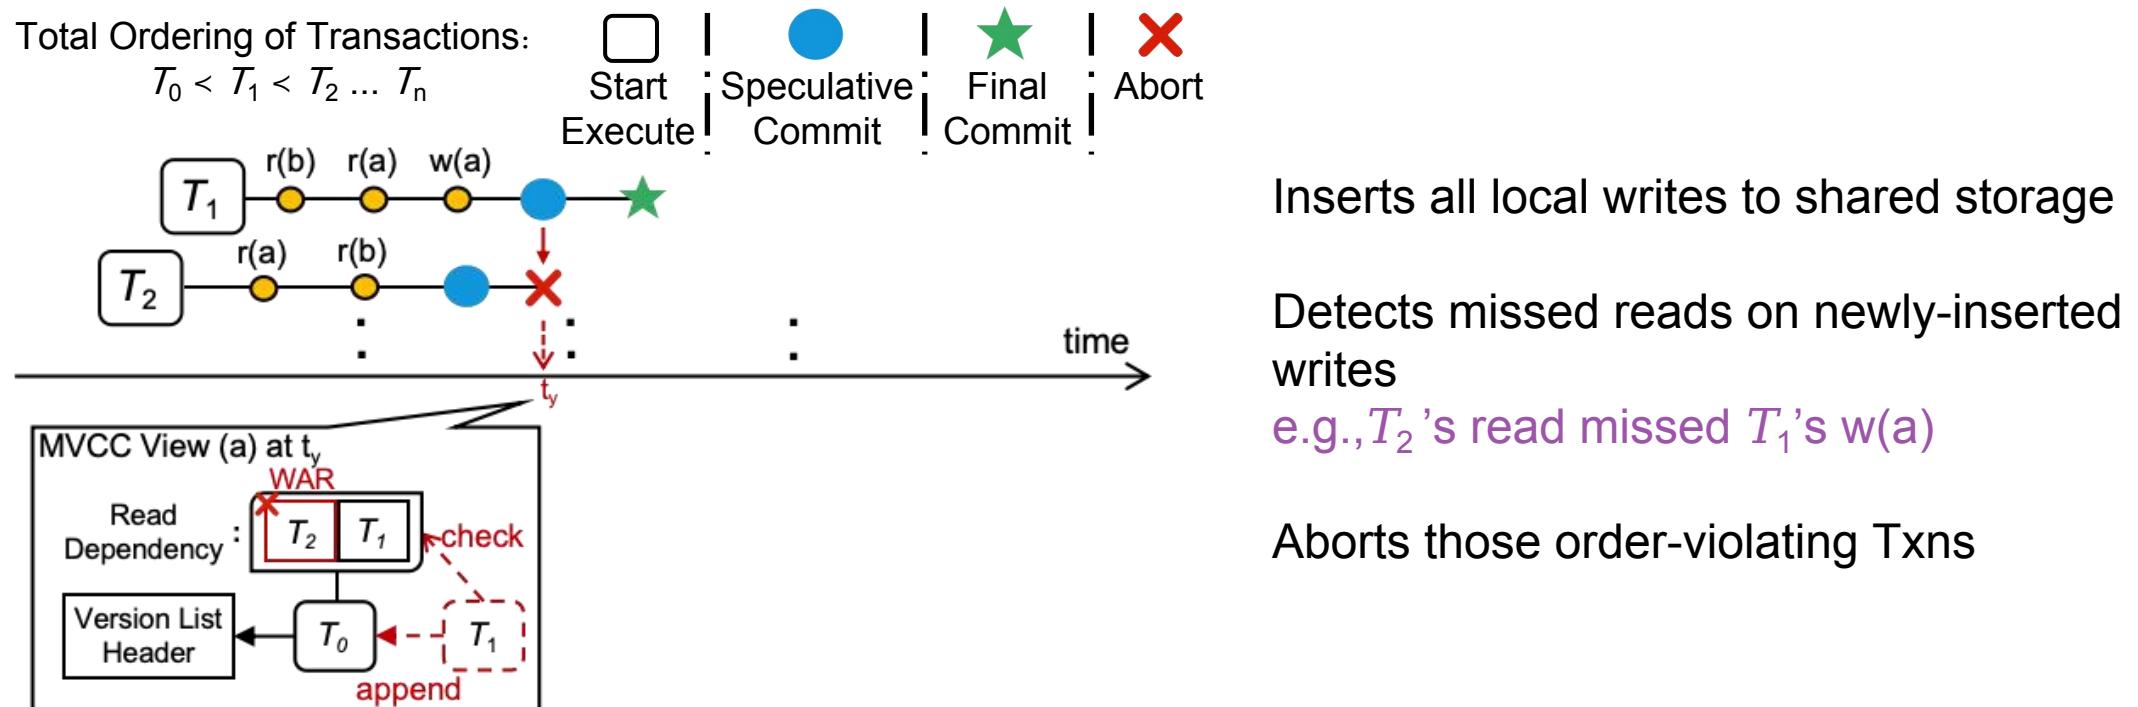


Methodology

■ #1 Speculative transaction execution

- Transaction lifecycle: 1) Execution; 2) Speculative Commit; 3) Final Commit

2) Speculative Commit: makes **writes visible** & **detects** conflicts & **aborts** order-violating Txns



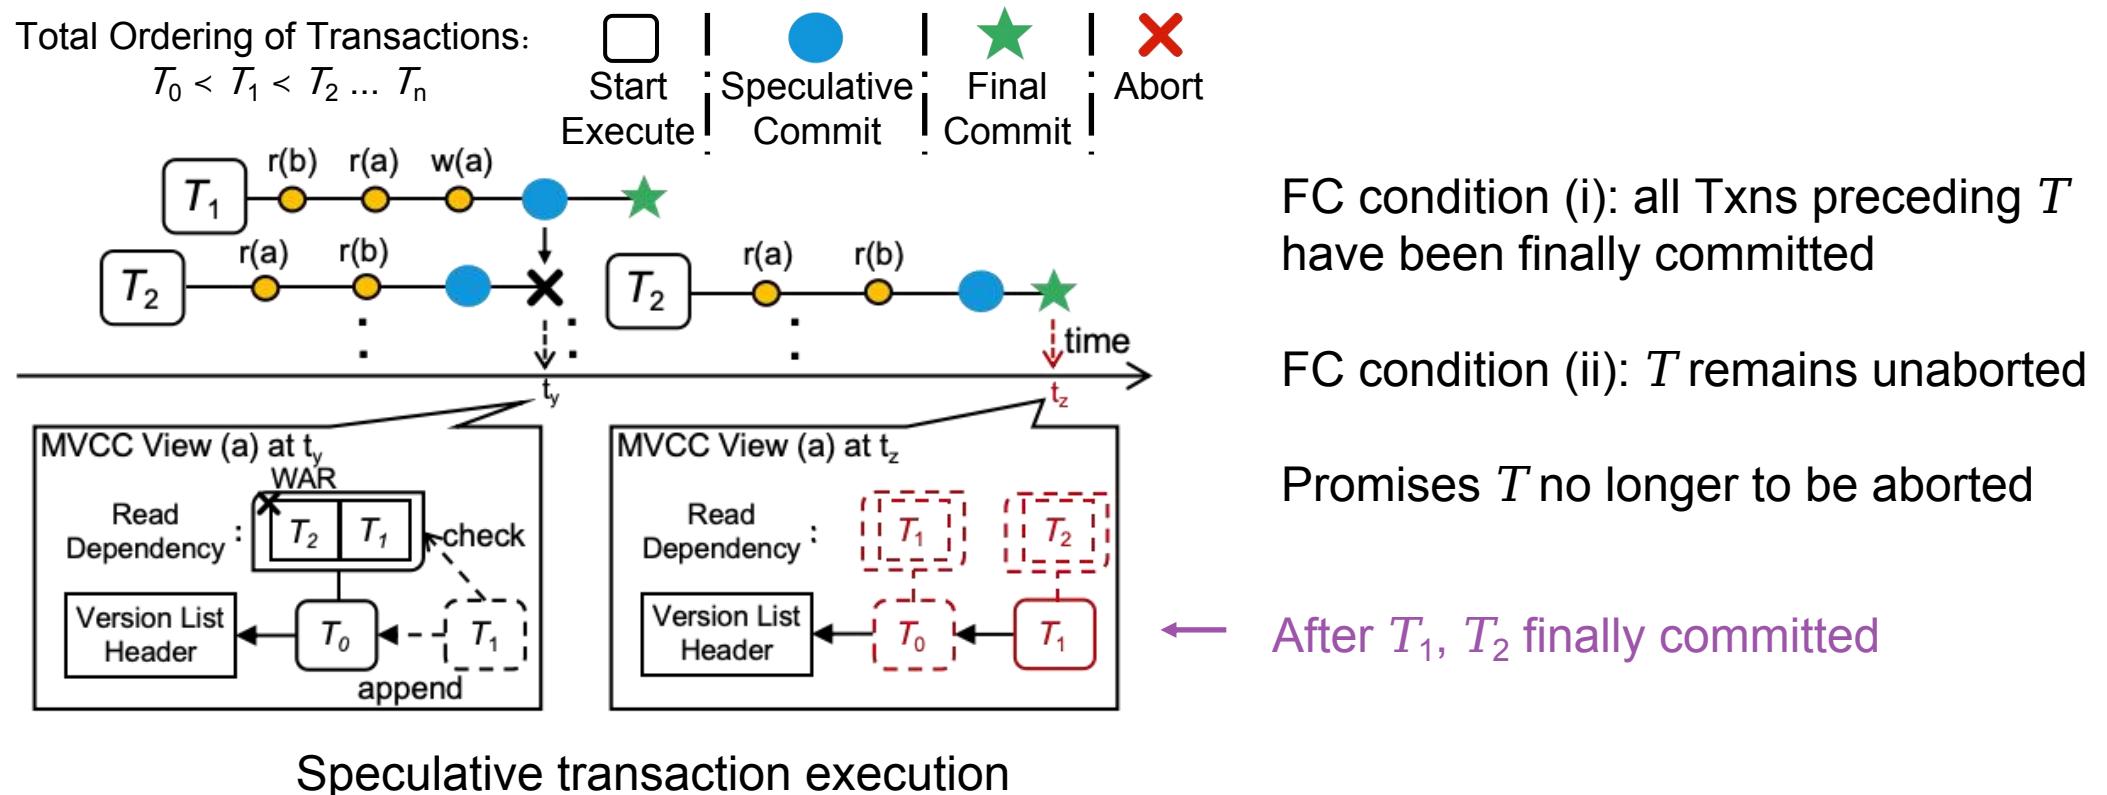
Speculative transaction execution

Methodology

#1 Speculative transaction execution

- Transaction lifecycle: 1) Execution; 2) Speculative Commit; 3) Final Commit

3) **Final Commit**: updates the tracking counter, **cleans** redundant versions & deps

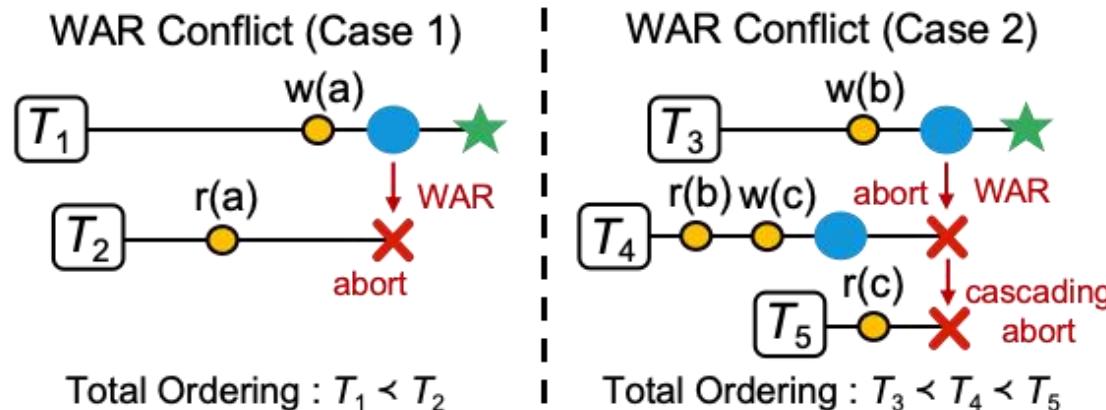


Methodology

Conflict types

- Mis-speculations lead to **write-after-read (WAR) conflict** (but no WAW or RAW conflicts)
- WAR conflict: Txn T_i 's write W_i (visible til T_i 's spec. commit) is **missed** by T_j 's read R_j , where W_i conflicts with R_j and $i < j$
- Become scaling bottlenecks under contended workloads

both the **Overhead** and the **Number** of mis-speculations



Exemplifying the WAR Conflict

Methodology

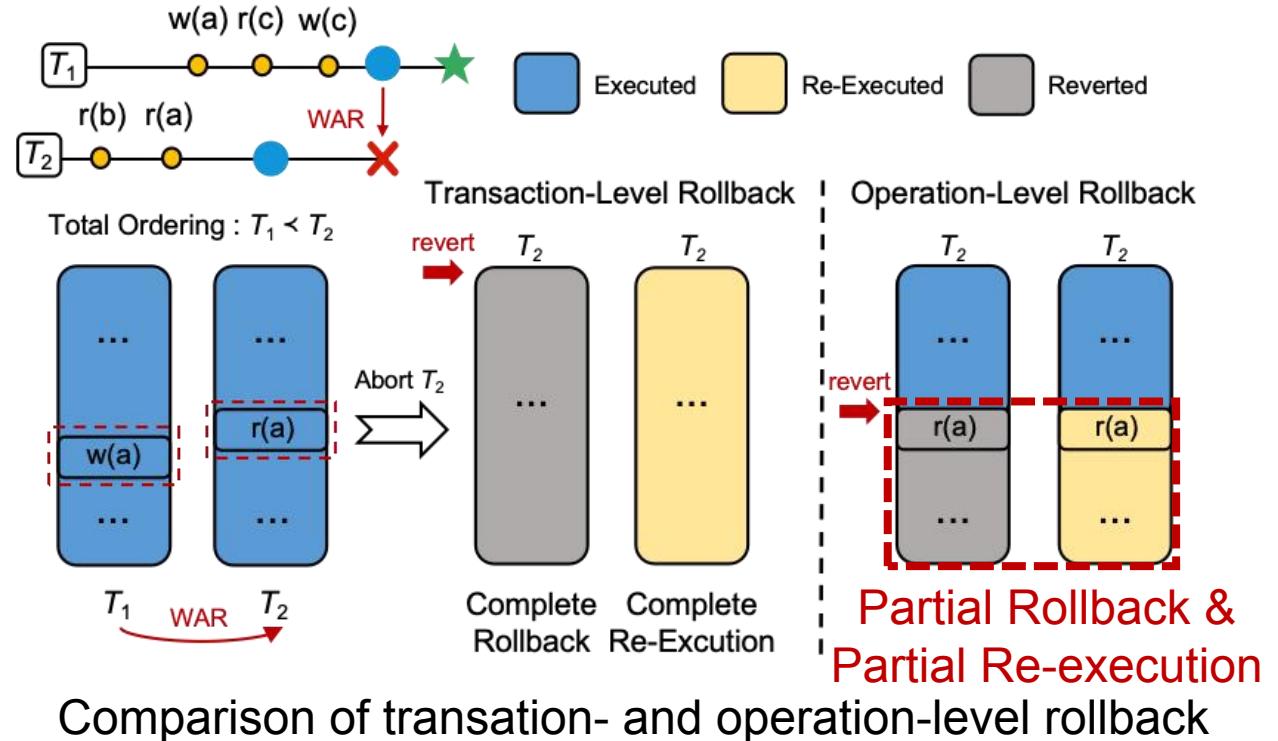
■ #2 Fine-Grained transaction rollback

- When a WAR conflict occurs, the aborted transaction(s) need **rollback** and **re-execute**
Expensive and wasteful for smart contract Txns !!!
- Partial rollback mechanism:
Only the operations impacted by the conflict
need to be roll-backed and re-executed
- Avoids wasting CPU resources
& Saves re-execution overhead
Reduces the overhead per mis-speculation

Traditionally: Complete
↓

Complete
↓

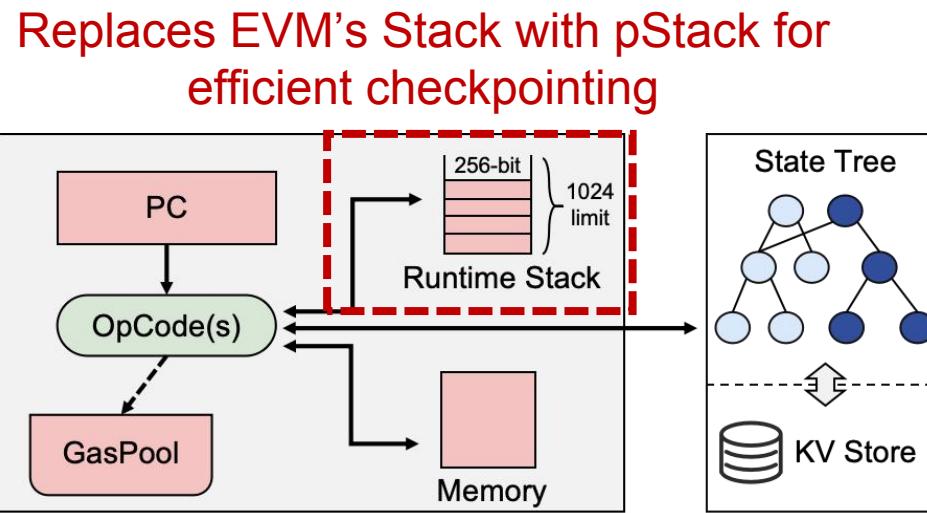
Expensive and wasteful for smart contract Txns !!!



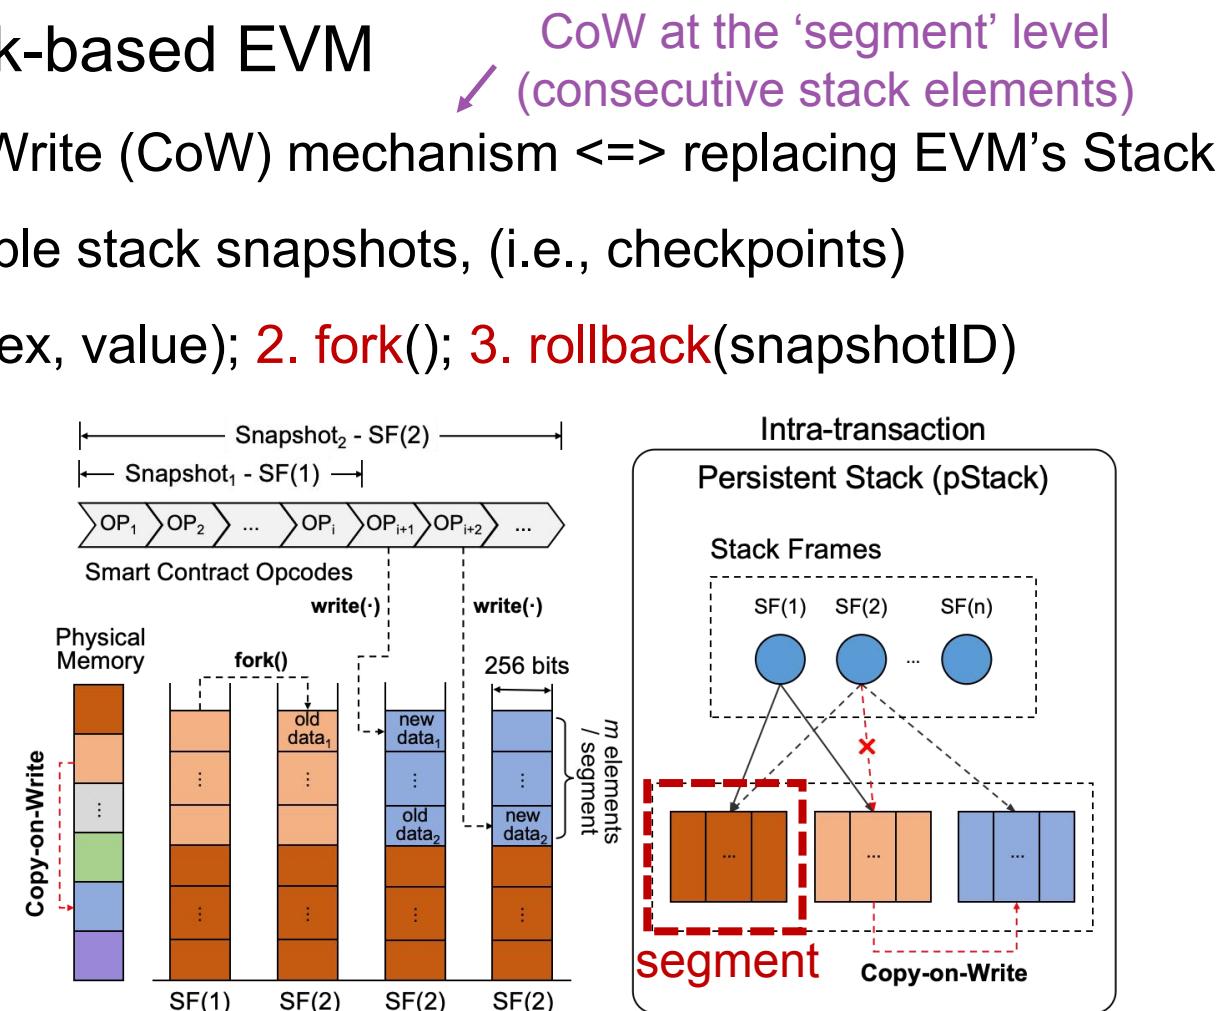
Methodology

- Implementing partial rollback in stack-based EVM
 - pStack, a persistent stack with Copy-on-Write (CoW) mechanism <=> replacing EVM's Stack
 - Efficiently restores and operates on multiple stack snapshots, (i.e., checkpoints)

Exposes three primitives: 1. `write(elementIndex, value)`; 2. `fork()`; 3. `rollback(snapshotID)`



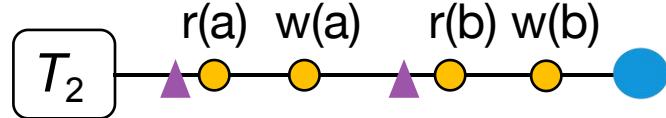
Ethereum Virtual Machine (EVM) structure



pStack-based stack frame management

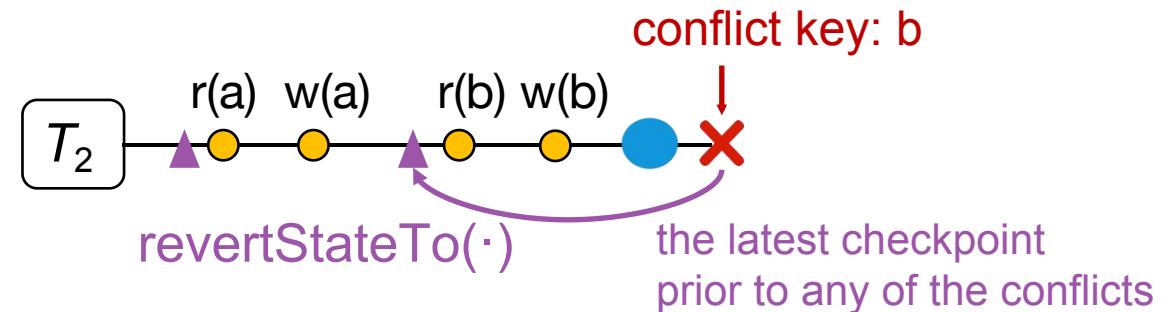
Methodology

- Integrates partial rollback implementation with speculative execution
- Makes a checkpoint **before an external read**, i.e., SLOAD with the read not in local storage
- Rollbacks to a checkpointed state **right before** the conflict occurs
- **Low checkpoint memory cost & cache-friendly**



▲ `makeCheckpoint(·)`:

- EVM.pStack.fork()
- Records PC, Memory
- EVM.addCheckpoint(·)

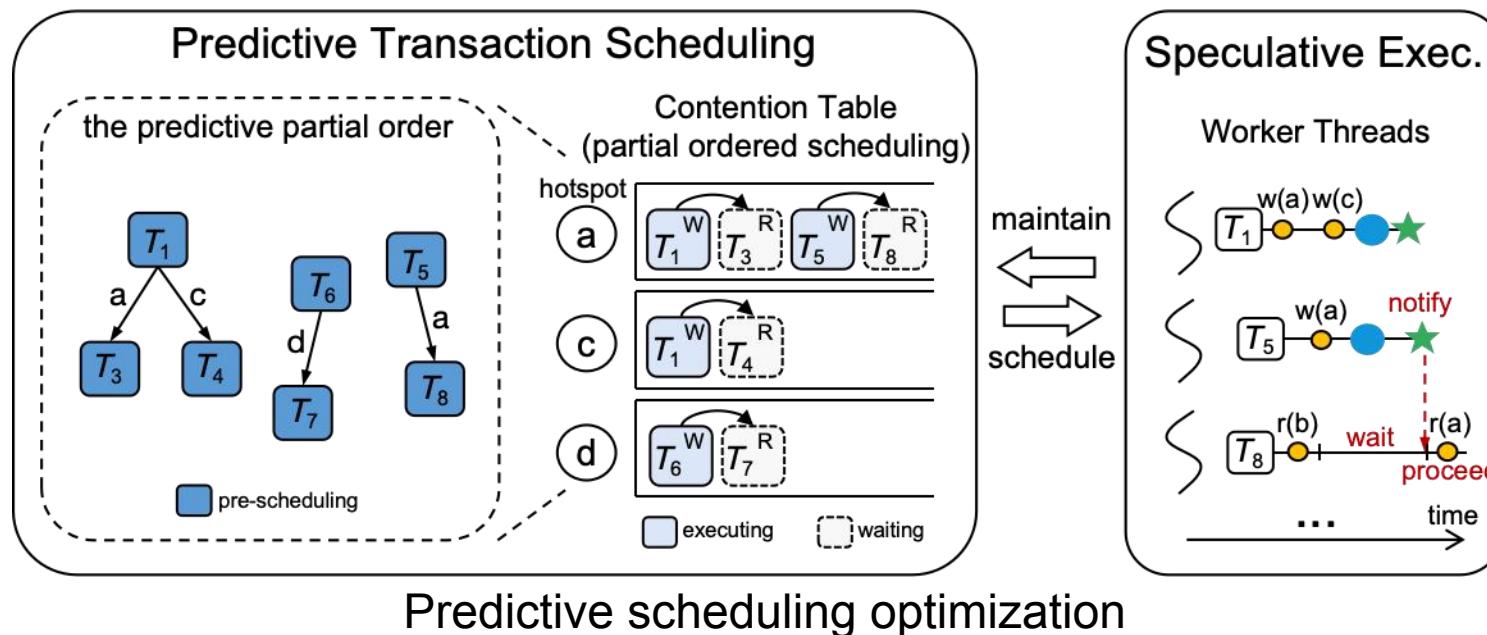


- EVM.getCheckpoint(·)
- Sets PC, Memory
- EVM.pStack.rollback(·)

Methodology

■ #3 Predictive transaction scheduling serving as ‘hints’

- Only requires **partial** a-priori read/write sets to pre-schedule potentially conflicting transactions
- If the prediction is correct, it effectively **reduces the number of runtime mis-speculations**
- **Independent** scheduling across replicas (ensures SDS despite different predictions)



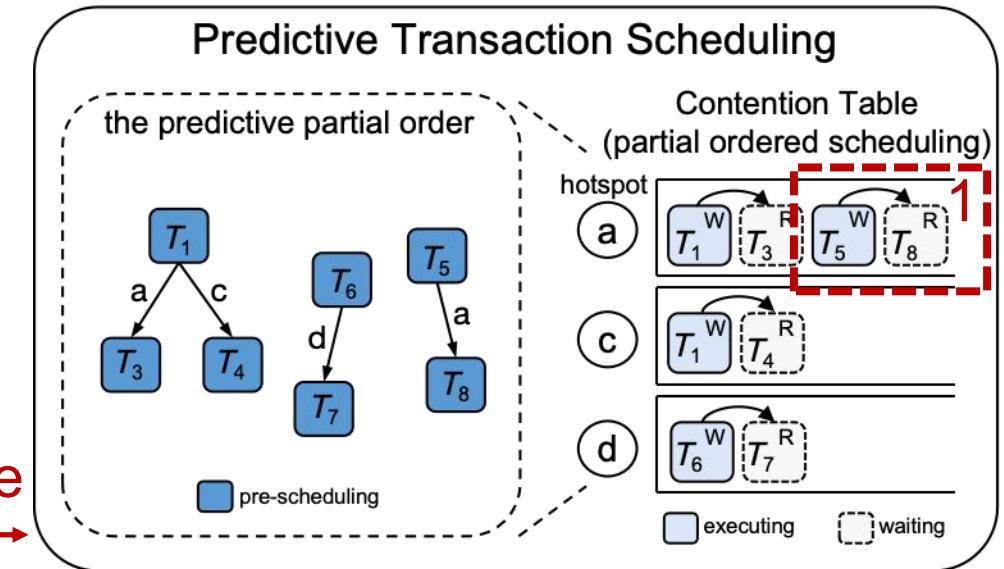
Methodology

- Key intuition: wait for preceding conflicting Txns to be finally committed => avoid WAR conflicts
- Sched. Algorithm (efficiency and correctness): concurrent maintenance and sequential granting
- Schedules **only highly-conflicted keys and predicted WAR conflicts**
- Superior parallelism than Calvin's single-threaded ordered lock pre-scheduling

Example:
pre-acquired rw_set

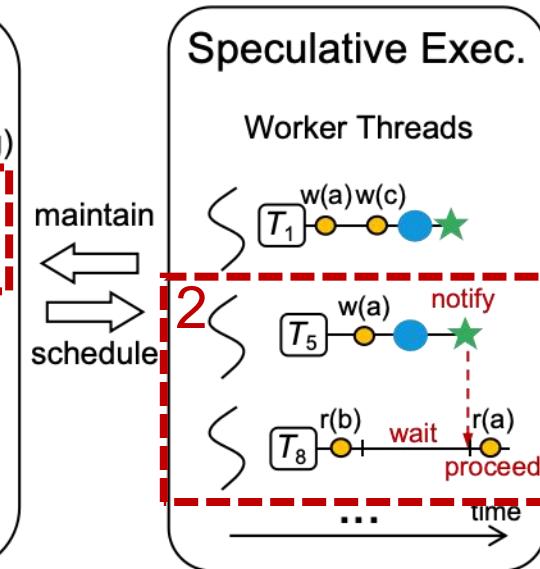
$T_5 : w(a)$
 $T_8 : r(a)$

1. Pre-maintenance



Predictive scheduling optimization

2. Runtime scheduling



Schedules reads, finer than Txns

T_8 's $r(a)$ waits for T_5 's final commit

Evaluation

■ Testbed

- 2x Intel Xeon Gold 6330 CPU (28C56T each) with 256GB DRAM

■ Benchmarks

- YCSB & SmallBank & TPC-C alike smart contracts (written in Solidity)
- Uniform workload and Skewed workload (controlled by *Zipf*)

■ Baselines (incl. SOTA DCC schemes)

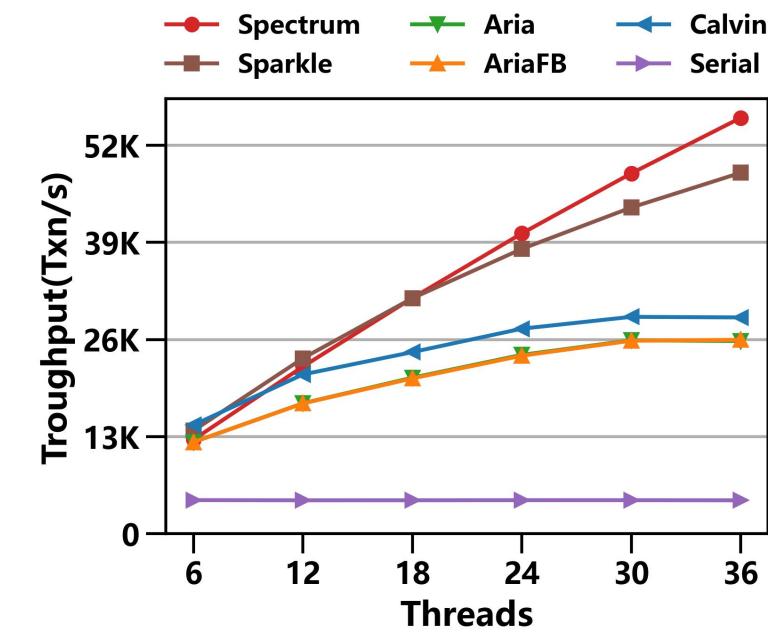
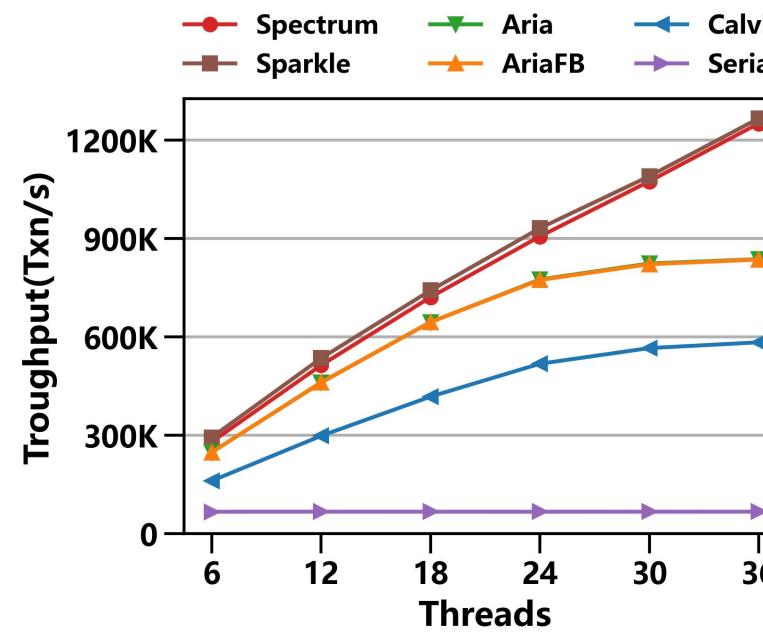
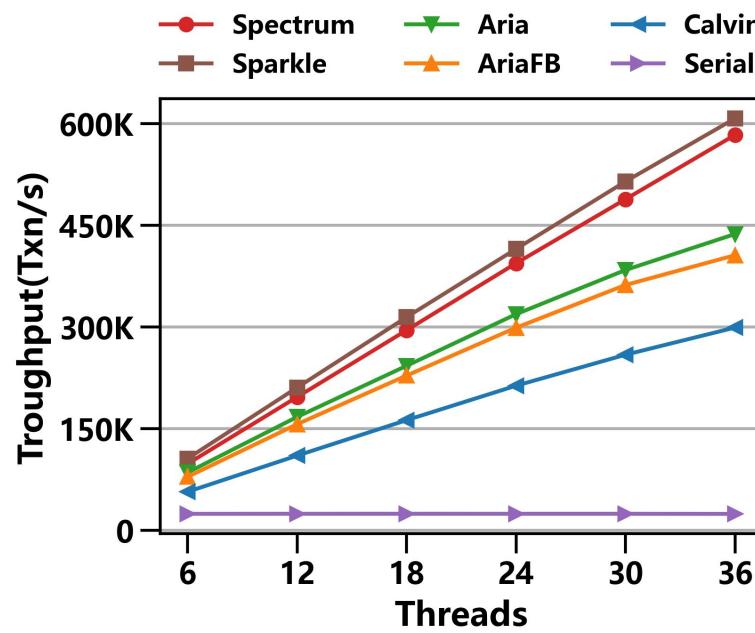
- Serial
- Calvin
- Aria, AriaFB
- Sparkle

Benchmark	Transaction Type (# of reads/writes)	call ratio
YCSB	Get&Set (5r5w)	100%
SmallBank	Six standard functions (2r0w, 2r1w, 1r1w * 2, 2r2w * 2)	equal probability
TPC-C	NewOrder, Delivery, Payment (4r5w/orderline, deliver 10 orderlines, 1r1w)	N:D:P = 11:11:1

Kindly refer to our paper for further impl. and experiments

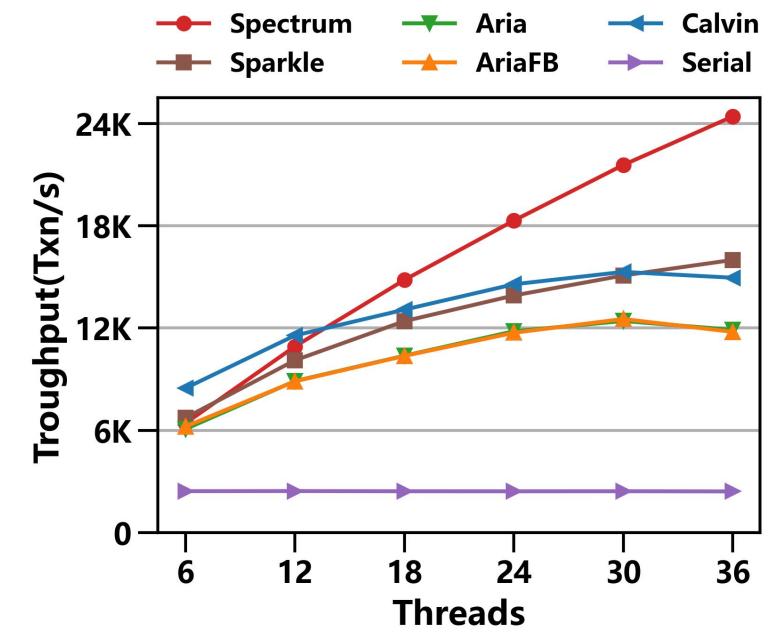
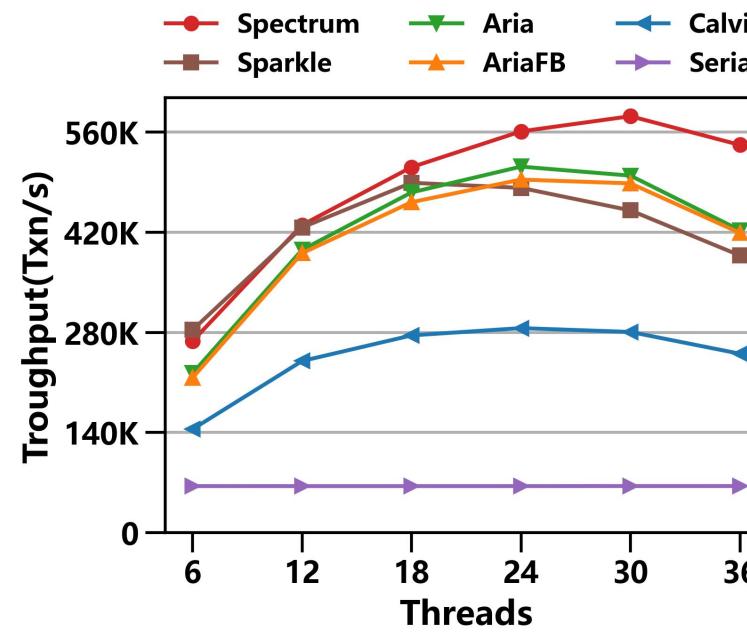
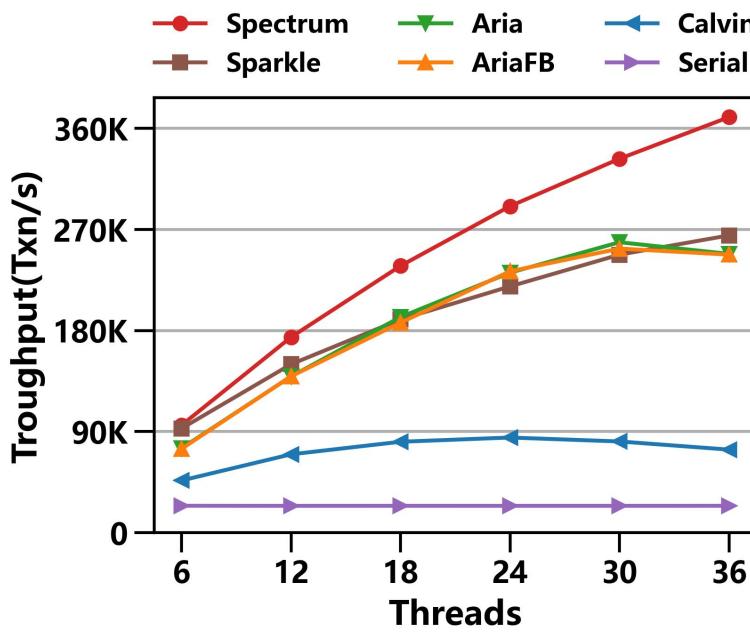
Evaluation

- Throughput of varying threads
 - Under the uniform workload, both Spectrum and Sparkle scale well within 36 threads, and outperform Aria (1.3x~2.1x), AriaFB (1.4x~2.1x), Calvin (1.7x~1.9x) and Serial (12.7x~24.2x)



Evaluation

- Throughput of varying threads
 - Under the skewed workload, Spectrum achieves the highest peak throughput
 - Up to 1.6x, 2.1x and 5.0x higher throughput than Sparkle, Aria, and Calvin

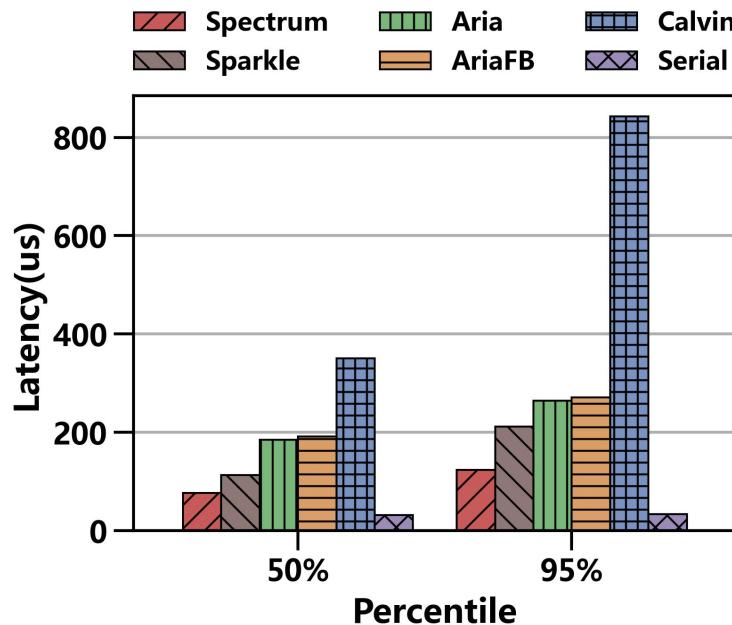


Evaluation

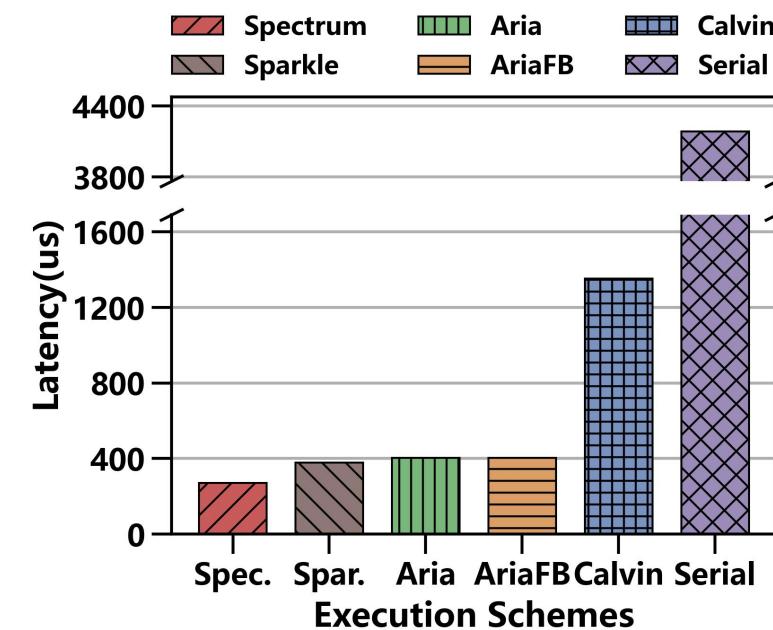
■

Transaction and Block latency

- Serial has the lowest tx latency, then Spectrum, Sparkle, Aria variants, and Calvin (the highest)
- Spectrum realizes the lowest block latency, **reducing it by 28.3%, 32.8%, 33.1%, 80.0%, and 93.5%** compared to Sparkle, Aria, AriaFB, Calvin, and Serial, respectively



p50, p95 Tx latency (YCSB, Zipf = 0.9)



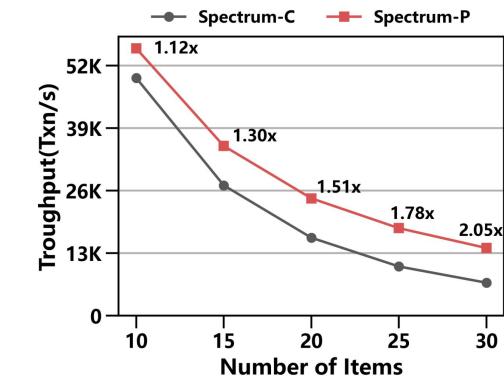
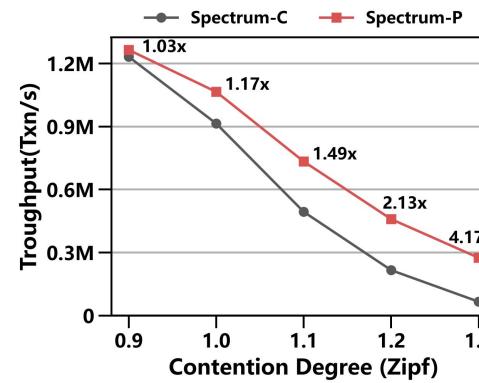
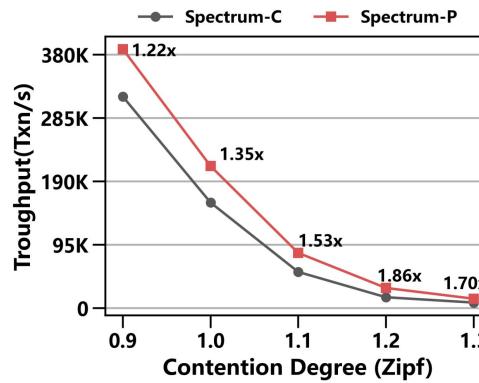
Block latency (YCSB, Zipf = 0.9)

Evaluation

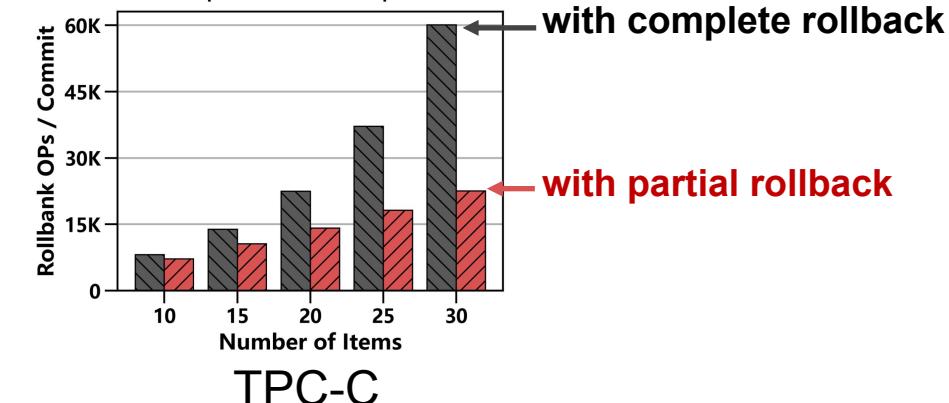
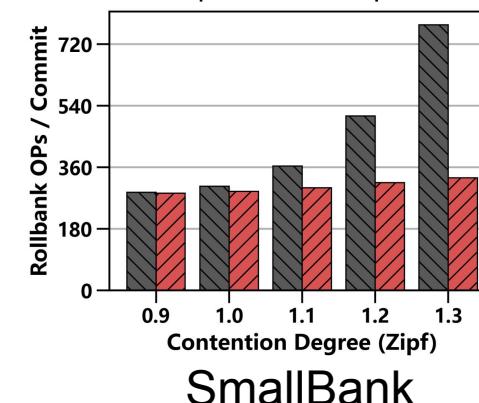
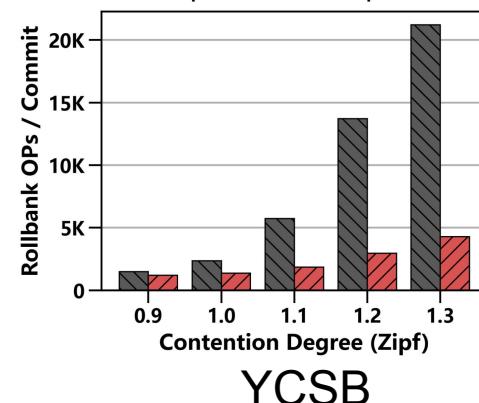
■ Evaluation of partial rollback (re-execution overhead)

- Partial rollback **effectively reduces** the number of operations to be rolled back and re-executed, by **79.8%**(YCSB), **57.6%**(Smallbank) and **64.9%**(TPC-C), compared to complete rollback

Scheme throughput
in varying Zipf



Rollback operations
per committed tx

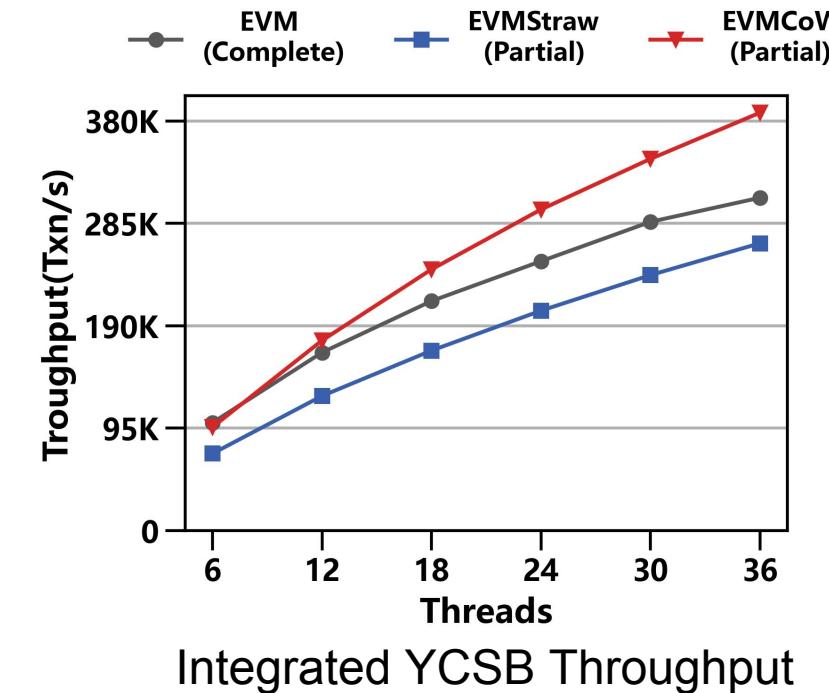
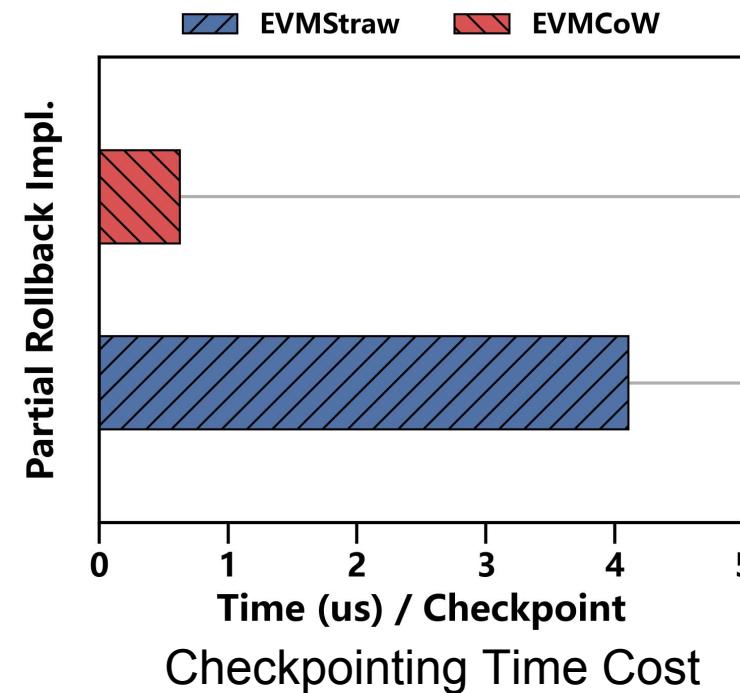
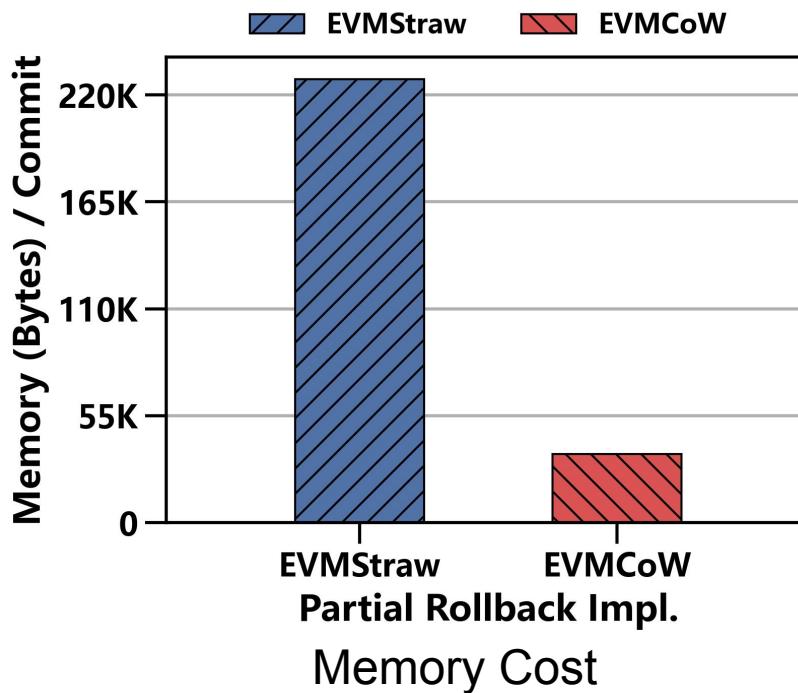


Evaluation

■ Evaluation of partial rollback (efficiency of pStack)

- pStack saves **84.8%** checkpointing memory cost and largely reduces creation time
- pStack enhances throughput by **31.2%** compared to the strawman approach
 - efficient impl. and cache-friendness

strawman: copies all stack elements when checkpointing
pStack: uses Copy-on-Write mechanism for checkpointing

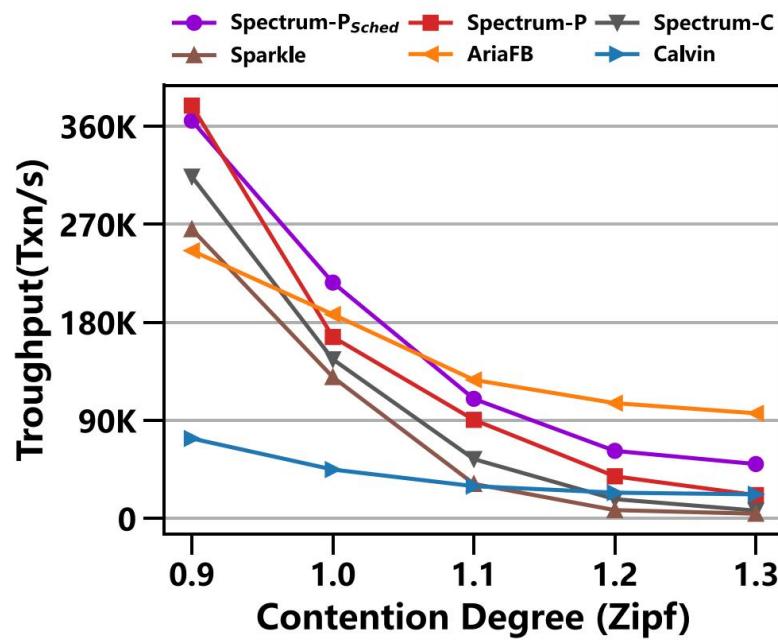


Evaluation

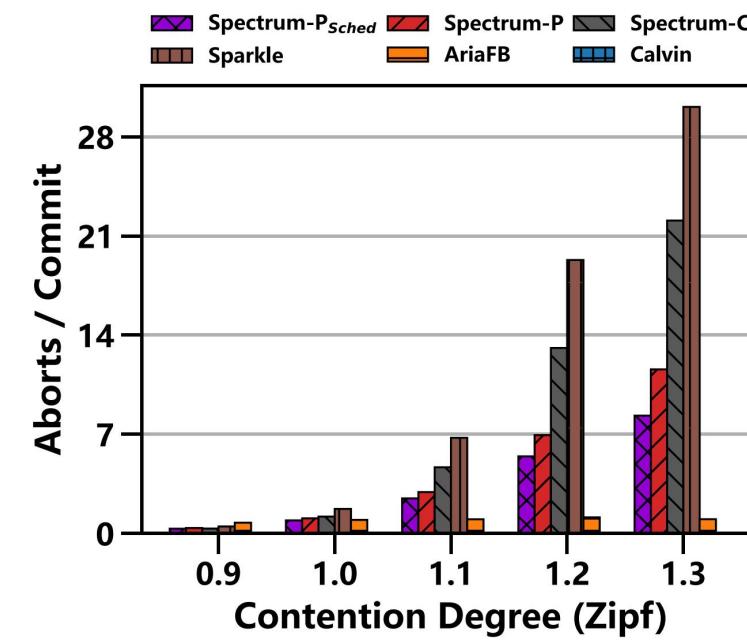
■ Evaluation of predictive scheduling

- Spectrum-P_{Sched} outperforms Spectrum-P by 2.3x and Spectrum-C by 6.8x
- Predictive scheduling further reduces aborts (number of mis-speculations) by 28.2% to 62.3%

Spectrum-P_{Sched} implements both predictive scheduling and partial rollback
Spectrum-P only applies partial rollback
Spectrum-C employs neither of these optimizations



Throughput of varying Zipf (YCSB)

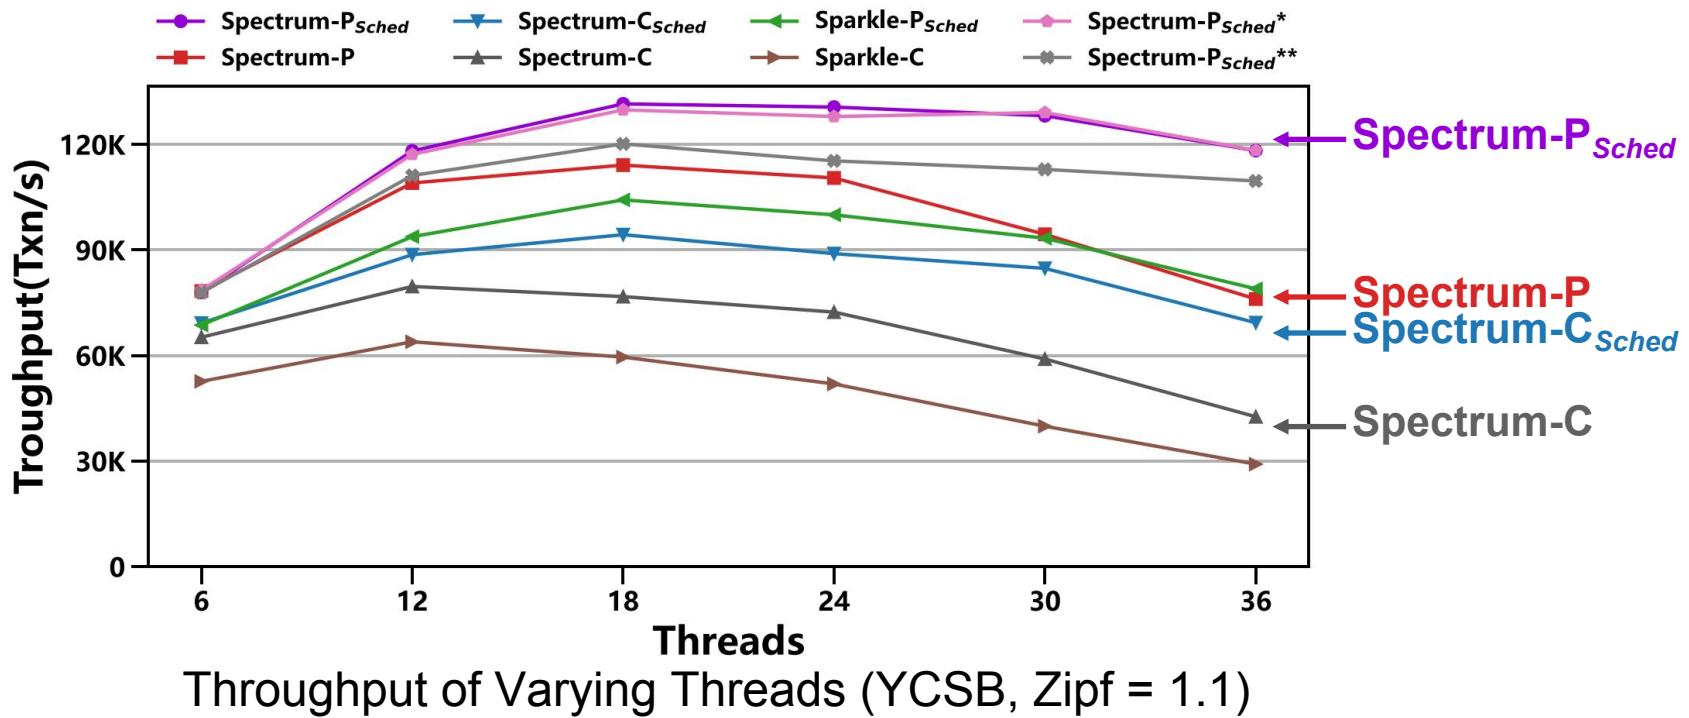


Aborts/Commit in Varying Zipf (YCSB)

Evaluation

■ Ablation Studies on Integrated Optimizations

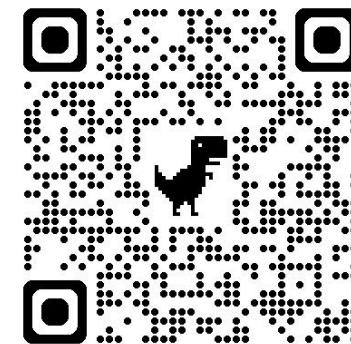
- The improvement of using only pre-scheduling (Spectrum-C_{Sched}, **1.6x**) is less than that of using only partial rollback (Spectrum-P, **1.8x**), but leveraging both yields the best performance (Spectrum-P_{Sched}, **2.8x** than Spectrum-C).



Conclusion & Future work

- Spectrum achieves both fair smart contract execution (by ensuring strict determinism) and high performance for blockchain ledgers
 - Speculative execution produces the same agreed-upon serial order with superior parallelism
 - Two novel optimizations: operation-level rollback & predictive scheduling
 - 1.4x ~ 4.1x higher throughput than SOTA DCC schemes (on YCSB, SmallBank, TPC-C alike smart contracts)
- Future work
 - Intra-transaction parallelism
 - Data-partitioned sharding settings

<https://github.com/jacklightChen/spectrum>



Paper Details



Source Code

Thank you for listening!

Q&A chenzh@stu.ecnu.edu.cn